



# DuckDB

An Embeddable Analytical Database

# Agenda

- **Motivation**
- Design
- Implementation
- Testing
- Next Steps

# DB is missing the boat again

- More and more complex Python/R/Julia/... libraries being deployed to solve basic relational problems
- DB world largely irrelevant there. Why?
  - Embeddability
  - Ease of Use

# What about SQLite ?

- In-process SQL database, data either in memory or in a file, rock-solid, used on every smartphone, browser, OS, ....
- People also use it for large-ish dataset analysis
- Bad idea, SQLite was never built for this
  - e.g. row-based storage model

# What about MonetDBLite?

- Attempted to re-design existing system. Sort of worked.
- Problems:
  - Error handling, global variables, restart, multi-DB, ...
  - Memory management & resource allocation difficult
  - Problematic processing paradigm for embedded
  - Bulk intermediates require lots of memory and/or disk space and interfere with host
  - No graceful handling of out-of-memory situations

Embedded



TERADATA

Stand-Alone



OLTP

OLAP

# Agenda

- Motivation
- **Design**
- Implementation
- Testing

# DuckDB Goals

- Fast OLAP, reasonable OLTP
  - e.g. concurrent appends
- Fully and easily embeddable
  - No globals, no dependencies
  - Works gracefully in low or out-of-memory situations
- Stability (aspiring to match SQLite)
- Clean, readable, consistent and extensible code
  - Basis for future research projects
- Full-featured and in public use



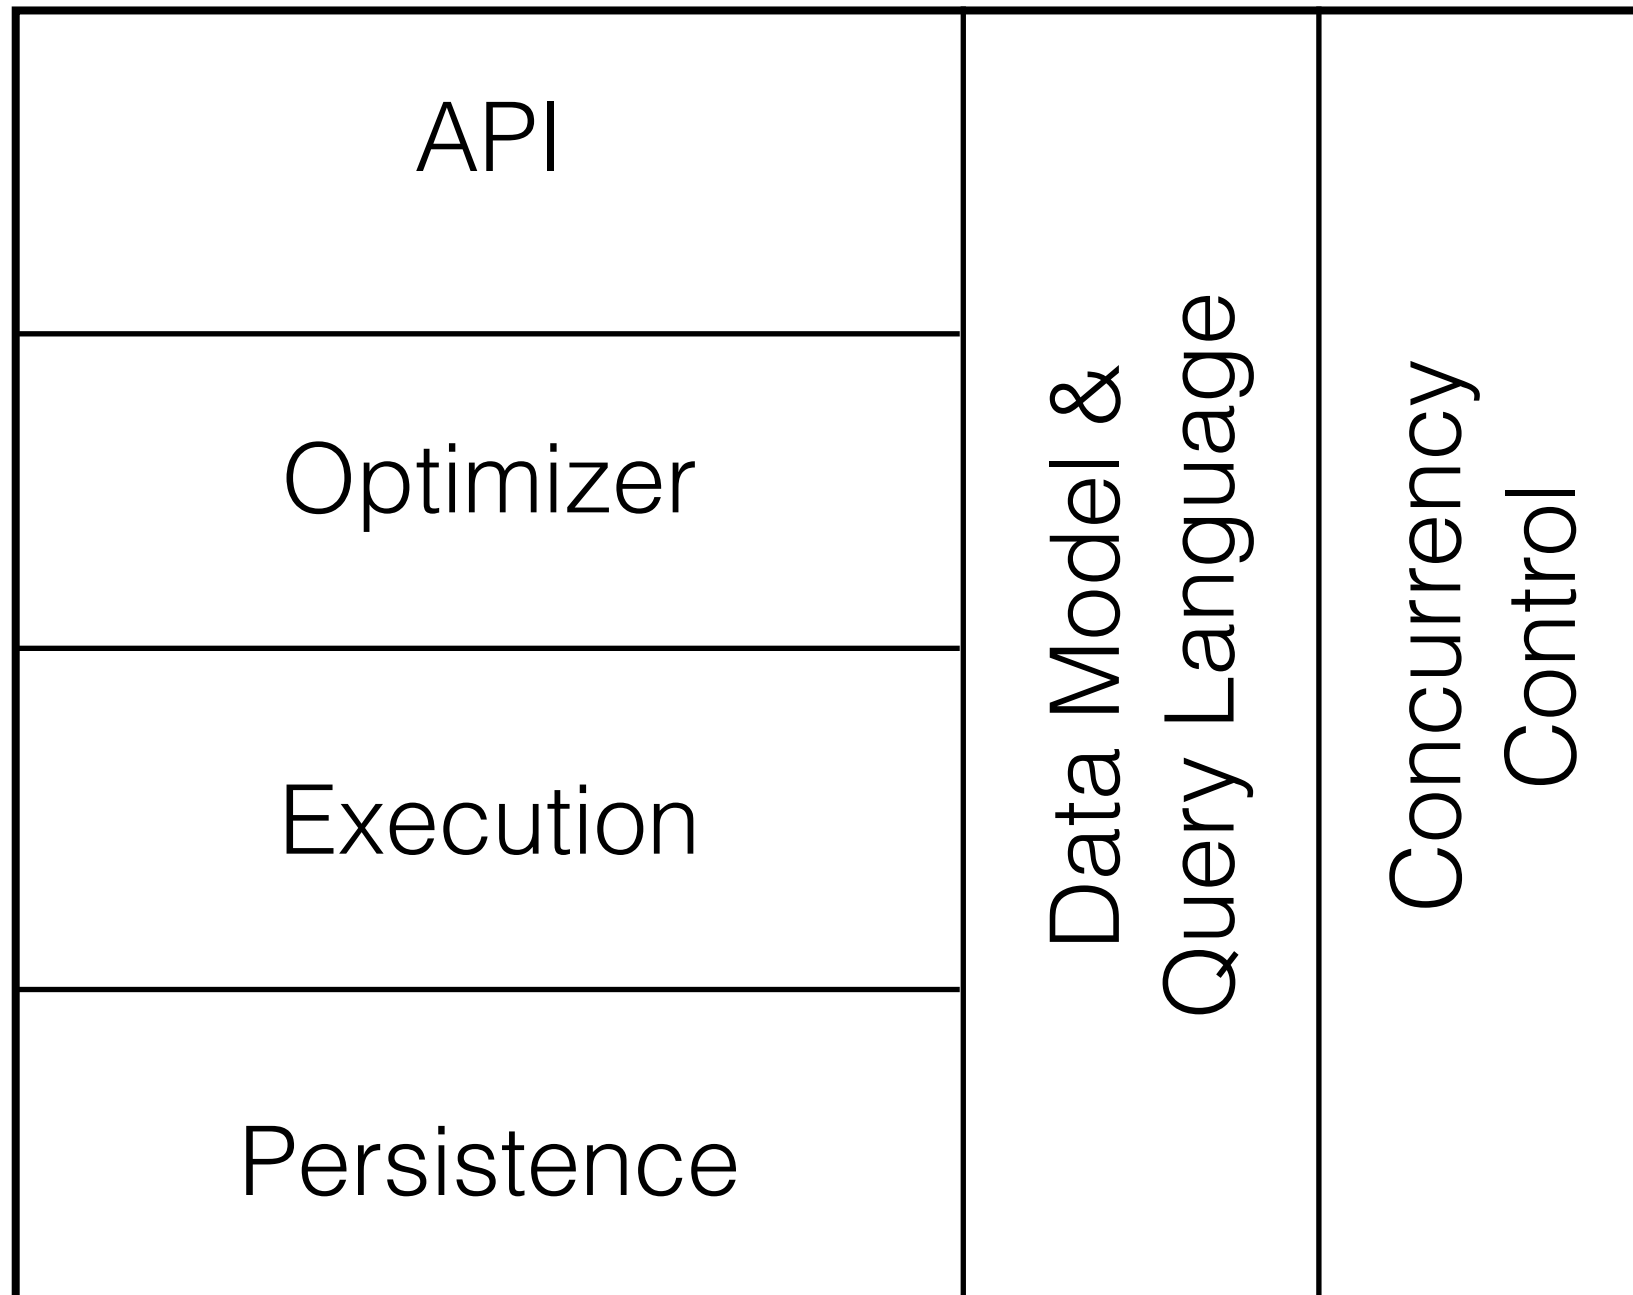
# DuckDB Design Choices

- SQL and relations
- Vectorized Model
  - because JIT has too many dependencies
- MVCC
  - e.g. concurrent appends

# Agenda

- Motivation
- Design
- **Implementation**
- Testing
- Next Steps

# Architecture Overview



# Implementation

- **C++ 11**
- Rich & efficient and stable standard library
  - no need to roll own lists/hash
  - not used for actual data but for all auxiliary structures
- Scopes do memory/lock management for us
  - unique\_ptrs, avoid leaks
  - lock\_guard, auto-acquiring and releasing locks
  - destructors
- OOP, information hiding, clean namespaces
  - important for embedding

# C++ STL?

- But “STL is slower than specialized solutions”
  - This is true.
- Specialized solutions for everything are prime example of premature optimization
- Without STL you need to do this for everything
- Often results in slower and buggier code
  - STL is extremely well tested (used by millions)

# Exceptions

- Exceptions + smart pointers = automatic cleanup
  - No leaks!
- Exceptions are zero-cost when not triggered
- In DuckDB: Never more than one per query. Ex:
  - Handle OOM
  - Query cancellation
  - Fatal query errors

# Exceptions

```
// update the statistics with the new data  
lock_guard<mutex> stats_lock(statistics_locks[column_id]);  
statistics[column_id].Update(updates.data[j]);
```

- Statistics::Update can throw an exception
- The lock will be released regardless, without us doing any cleanup

# Templates

- Code expansion in C: scripts/macros
- Code expansion in C++: Templates
  - Debuggable, more readable, less error-prone



# Life of a query

- ```
SELECT count(*)  
FROM lineitem JOIN orders ON  
l_orderkey=o_orderkey  
WHERE o_orderstatus='X' AND l_tax > 50;
```



# Step 1: Parser

- PostgreSQL parser
  - Battle-tested
  - Side-effect: Postgres compatibility
  - `libpg_query`
- Transform into custom class structure
  - Inspired by Peloton



```

//! SelectStatement is a typical SELECT clause
class SelectStatement : public SQLStatement {
    public:
        SelectStatement()
            : SQLStatement(StatementType::SELECT), select_distinct(false),
              union_select(nullptr){};

        //! The projection list
        vector<unique_ptr<Expression>> select_list;
        //! The FROM clause
        unique_ptr<TableRef> from_table;
        //! The WHERE clause
        unique_ptr<Expression> where_clause;
        //! DISTINCT or not
        bool select_distinct;
        //! Group By Description
        GroupByDescription groupby;
        //! Order By Description
        OrderByDescription orderby;
        //! Limit Description
        LimitDescription limit;

        unique_ptr<SelectStatement> union_select;
        unique_ptr<SelectStatement> except_select;
}

```



# Step 2: Binder & Planner

- Binder
  - Resolve table and column names
  - Resolve data types
  - Overflow prevention!
    - Data statistics used for type promotion if required
    - Statistics right now: min, max, max str length
- Planner
  - Transform parse tree into logical operator tree



```
SELECT count(*)  
FROM lineitem JOIN orders ON l_orderkey=o_orderkey  
WHERE o_orderstatus='X' AND l_tax > 50;
```



```
AGGREGATE_AND_GROUP_BY[COUNT_STAR]  
  FILTER[(l_tax > CAST[DECIMAL](50)), (o_orderstatus = X)]  
  JOIN[EQUAL(l_orderkey, o_orderkey)]  
    GET(lineitem)  
    GET(orders)
```



# Step 3: Optimizer

- Rule-based optimizer
  - Matches tree patterns
- No join ordering yet

```
ConstantFoldingRule::ConstantFoldingRule() {  
    root = std::unique_ptr<AbstractRuleNode>(new ExpressionNodeSet(  
        {ExpressionType::OPERATOR_ADD, ExpressionType::OPERATOR_SUBTRACT,  
          ExpressionType::OPERATOR_MULTIPLY, ExpressionType::OPERATOR_DIVIDE,  
          ExpressionType::OPERATOR_MOD}));  
    root->children.push_back(  
        make_unique_base<AbstractRuleNode, ExpressionNodeType>(  
            ExpressionType::VALUE_CONSTANT));  
    root->children.push_back(  
        make_unique_base<AbstractRuleNode, ExpressionNodeAny>());  
    root->child_policy = ChildPolicy::UNORDERED;  
}
```

```

AGGREGATE_AND_GROUP_BY [COUNT_STAR]
  FILTER [ (l_tax > CAST[DECIMAL] (50)) , (o_orderstatus = X) ]
  JOIN [EQUAL (l_orderkey, o_orderkey) ]
  GET (lineitem)
  GET (orders)

```



```

AGGREGATE_AND_GROUP_BY [COUNT_STAR]
  JOIN [EQUAL (l_orderkey, o_orderkey) ]
  FILTER [ (l_tax > 50.000000) ]
  GET (lineitem)
  FILTER [ (o_orderstatus = X) ]
  GET (orders)

```

Pushdown!



# Step 4: Physical Planner

- Selects physical implementation for logical operators

```
void PhysicalPlanGenerator::Visit(LogicalJoin &op) {
    if (has_equality) {
        // equality join: use hash join
        plan = make_unique<PhysicalHashJoin>(move(left), move(right),
  move(op.conditions), op.type);
    } else {
        // non-equality join: use nested loop
        if (op.type == JoinType::INNER) {
            plan = make_unique<PhysicalNestedLoopJoinInner>(
                move(left), move(right), move(op.conditions), op.type);
        } else if (op.type == JoinType::ANTI || op.type == JoinType::SEMI) {
            plan = make_unique<PhysicalNestedLoopJoinSemi>(
                move(left), move(right), move(op.conditions), op.type);
        } else {
            throw NotImplementedException(
                "Unimplemented nested loop join type!");
        }
    }
}
```



```
AGGREGATE_AND_GROUP_BY[COUNT_STAR]  
  JOIN[EQUAL(l_orderkey, o_orderkey)]  
    FILTER[(l_tax > 50.000000)]  
      GET(lineitem)  
    FILTER[(o_orderstatus = X)]  
      GET(orders)
```



```
HASH_GROUP_BY[COUNT_STAR]  
  HASH_JOIN[EQUAL(l_orderkey, o_orderkey)]  
    FILTER[(l_tax > 50.000000)]  
      SEQ_SCAN[lineitem]  
    FILTER[(o_orderstatus = X)]  
      SEQ_SCAN[orders]
```



# Step 5: Execution

- **DataChunk** with max. length 1024 (Table slice)
  - **Vectors**, which are native arrays of certain type (int, float etc.)
  - NULL masks (16 x 8 byte integers per vector!)
    - Can be inherited or OR-ed together for vector operations
  - Selection vectors
- “Vector-Volcano”: pull DataChunk from root node of plan
  - Continue until result is empty, query then finished
  - Early materialisation
- Physical operators implemented using library of vector operations



# Query Profiling

Never scanned  
because RHS empty



|                                                                 |                                                |
|-----------------------------------------------------------------|------------------------------------------------|
| HASH_GROUP_BY<br>(0.00s)<br>1                                   |                                                |
| HASH_JOIN<br>INNER<br>l_orderkey<br>=o_orderkey<br>(0.00s)<br>0 |                                                |
| FILTER<br>l_tax ><br>50.000000<br>(0.00s)<br>0                  | FILTER<br>o_orderstatus<br>= X<br>(0.00s)<br>0 |
| SEQ_SCAN<br>lineitem<br>(0.00s)<br>0                            | SEQ_SCAN<br>orders<br>(0.00s)<br>150000        |



# DataChunk & Vector

```
class Vector {  
    public:  
        //! The type of the elements stored in the vector.  
        TypeId type;  
        //! The amount of elements in the vector.  
        size_t count;  
        //! A pointer to the data.  
        char *data;  
        //! The selection vector of the vector.  
        sel_t *sel_vector;  
        //! The null mask of the vector, if the Vector has any NULL values  
        nullmask_t nullmask;  
}  
  
class DataChunk {  
    public:  
        //! The amount of vectors that are part of this DataChunk.  
        size_t column_count;  
        //! The vectors owned by the DataChunk.  
        std::unique_ptr<Vector[]> data;  
        //! The (optional) selection vector of the DataChunk. Each of the member  
        //! vectors reference this selection vector.  
        sel_t *sel_vector;  
}
```

# VectorOperations

```
void VectorOperations::Add(Vector &left, Vector &right, Vector &result) {  
    switch (left.type) {  
        ...  
        case TypeId::INTEGER:  
            templated_binary_loop<int32_t, operators::Add>(left, right, result);  
            break;  
        ...  
    }  
  
    struct Add {  
        template <class T> static inline T Operation(T left, T right) {  
            return left + right;  
        }  
    };  
};
```

```

template <class T, class OP>
void templated_binary_loop(Vector &left, Vector &right, Vector &result) {
    auto ldata = (T *)left.data;
    auto rdata = (T *)right.data;
    auto result_data = (T *)result.data;

    result.nullmask = left.nullmask | right.nullmask;
    binary_loop_function_array<T, OP>(
        ldata, rdata, result_data, left.count, left.sel_vector);

    result.sel_vector = left.sel_vector;
    result.count = left.count;
}

```

## Template Magic

```

template <class T, class OP>
static inline void
binary_loop_function_array(T *__restrict ldata,
                           T *__restrict rdata,
                           T *__restrict result_data, size_t count,
                           sel_t *__restrict sel_vector) {
    if (sel_vector) {
        for (size_t i=0; i < count; i++) {
            result_data[sel_vector[i]] = OP::Operation(ldata[sel_vector[i]], rdata[sel_vector[i]]);
        }
    } else {
        for (size_t i=0; i < count; i++) {
            result_data[i] = OP::Operation(ldata[i], rdata[i]);
        }
    }
}

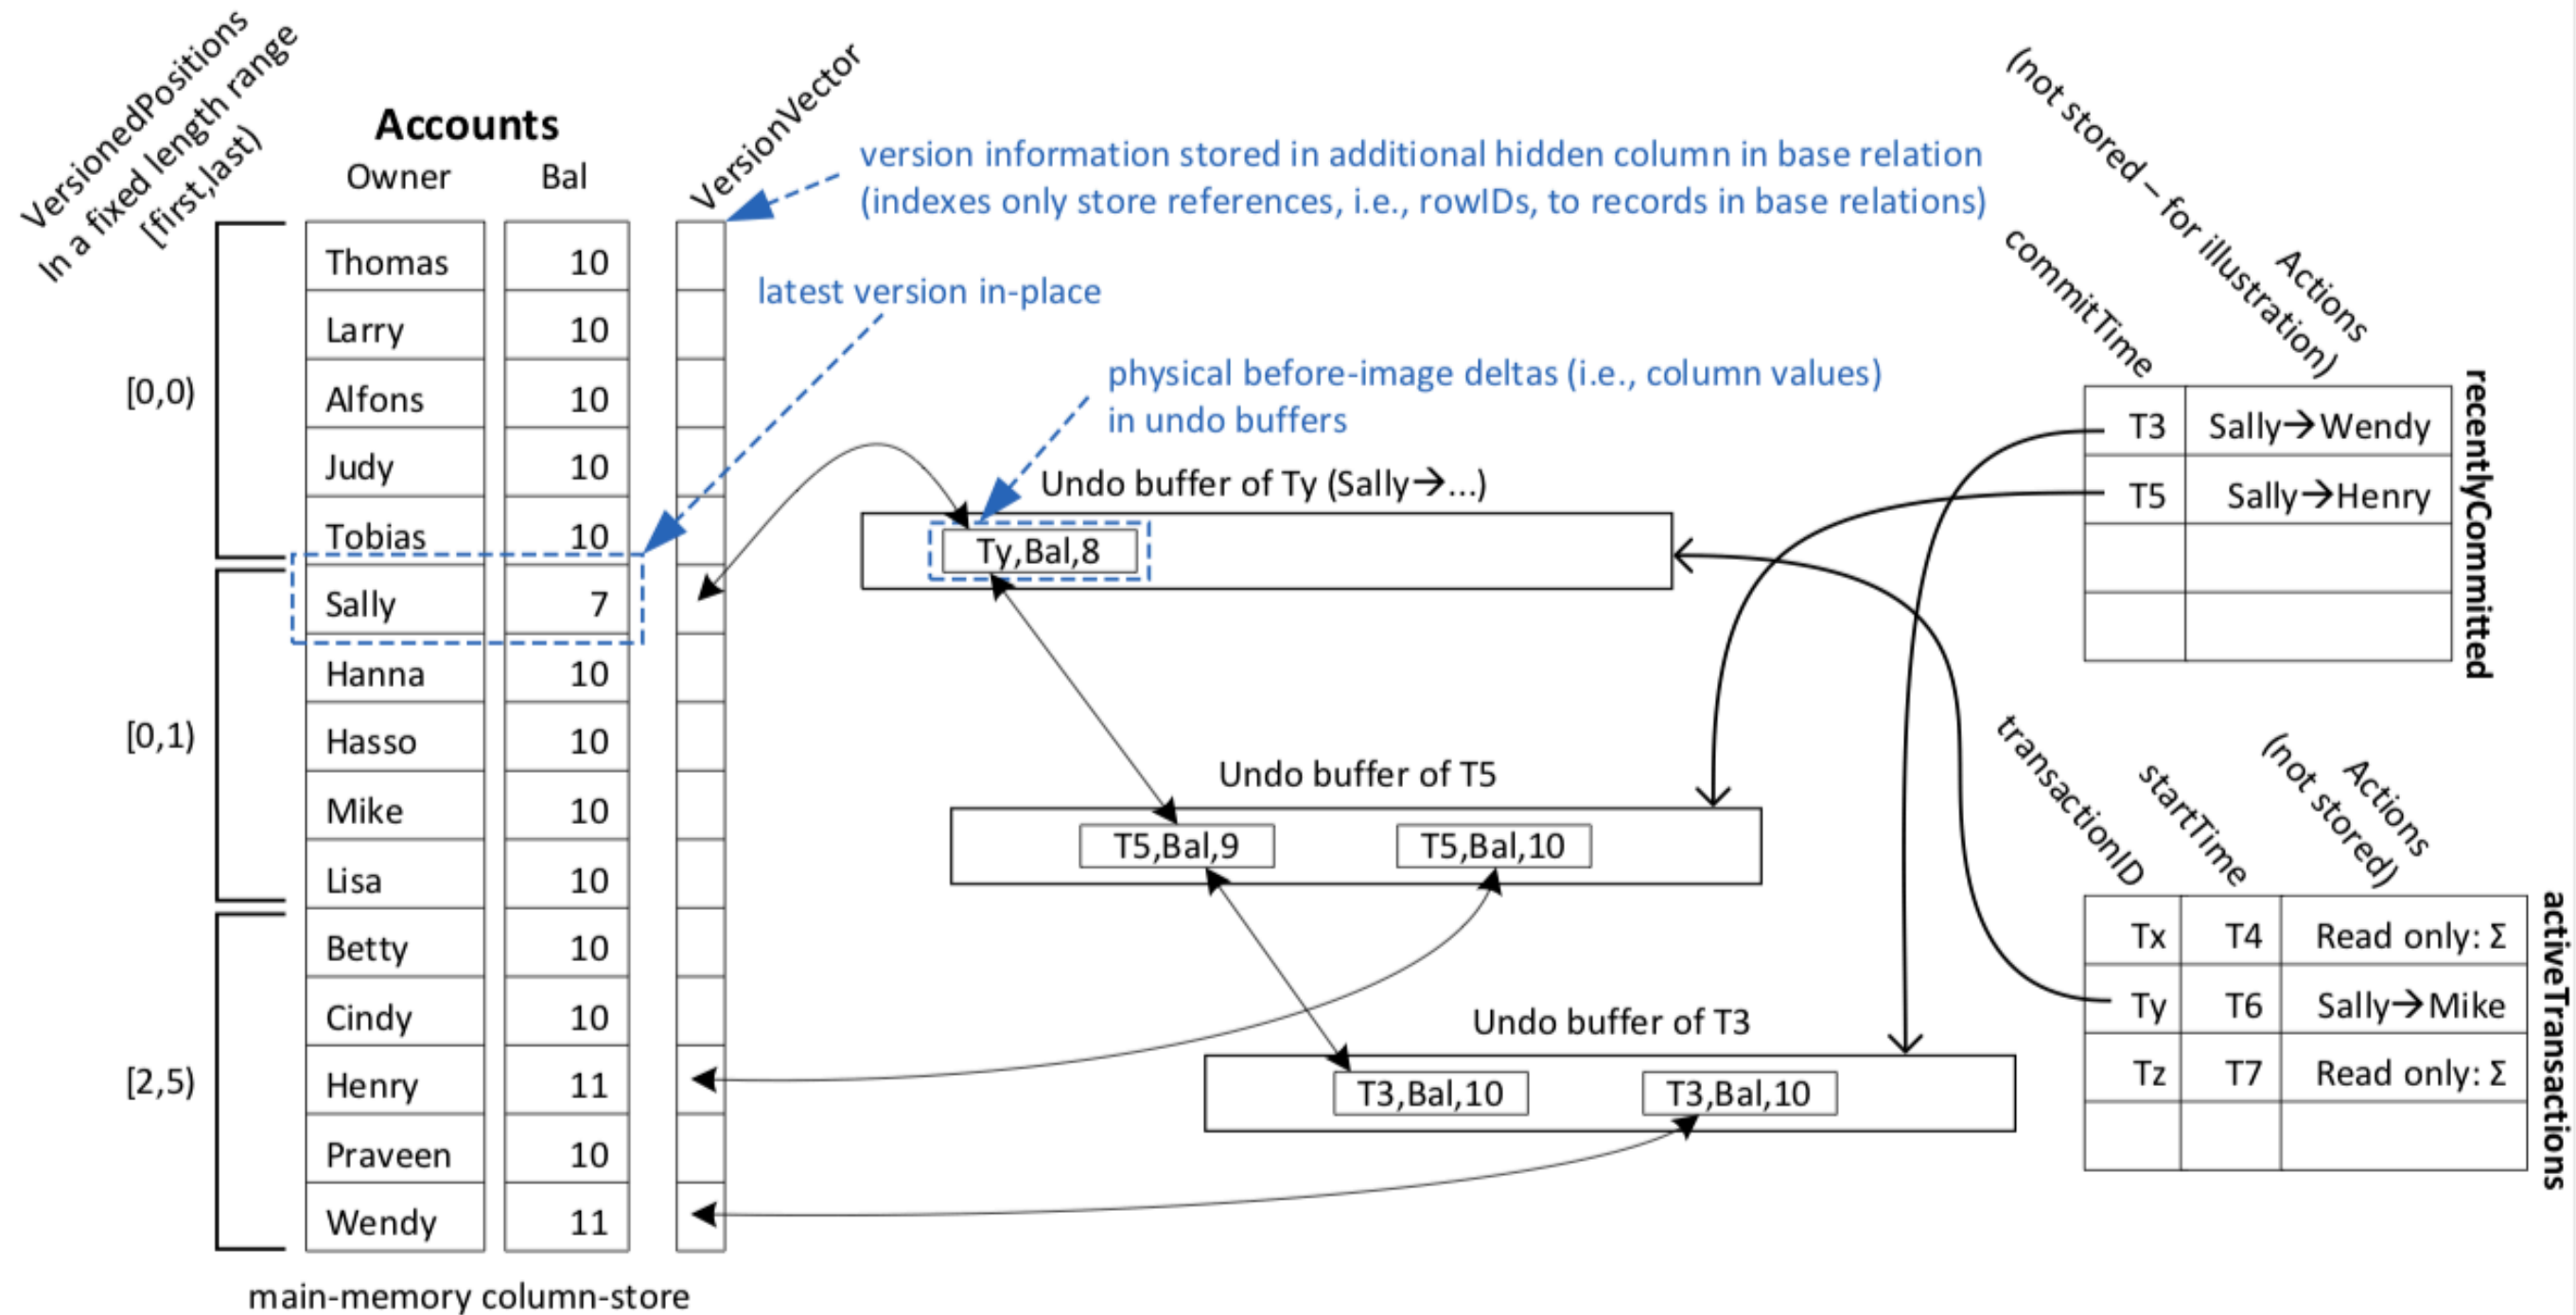
```

# Templating result:

```
for (size_t i = 0; i < count; i++) {  
    result_data[i] = ldata[i] + rdata[i];  
}
```

- Tight loops for every type
- Compiler will SIMDize loops
  - restrict in template

# MVCC Design

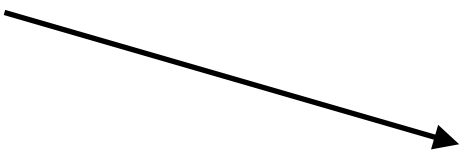


"Fast Serializable Multi-Version Concurrency Control for Main-Memory Database Systems"  
Thomas Neumann, Tobias Mühlbauer and Alfons Kemper.



# DuckDB API

- Main API: C++
- C API
  - duckdb\_open()
  - duckdb\_connect()
  - duckdb\_query()
  - ...
- SQLite API wrapper (same header)
- SQLite shell (demo!)



```
DuckDB db(nullptr);
DuckDBConnection con(db);

auto res = con.Query("SELECT 42");
int result = res->GetValue<int>(0, 0);
```

# Agenda

- Motivation
- Design
- Implementation
- **Testing**
- Next Steps

# Testing overview

- Unit tests with Catch2
  - core: ~one minute (demo)
  - extended: 30 mins
- SQLite sqllogictests
- sqlsmith
- Continuous integration with Jenkins
- Continuous code coverage with lcov



# Example Unit Test Case

```
#include "catch.hpp"
```

```
TEST_CASE("Test LEFT OUTER JOIN", "[join]") {  
    unique_ptr<DuckDBResult> result;  
    DuckDB db(nullptr);  
    DuckDBConnection con(db);  
  
    con.Query("CREATE TABLE integers(i INTEGER, j INTEGER)");  
    con.Query("INSERT INTO integers VALUES (1, 2), (2, 3), (3, 4)");  
    con.Query("CREATE TABLE integers2(k INTEGER, l INTEGER)");  
    con.Query("INSERT INTO integers2 VALUES (1, 10), (2, 20)");  
  
    result = con.Query("SELECT * FROM integers LEFT OUTER JOIN integers2 ON "  
                        "integers.i=integers2.k ORDER BY i");  
  
    REQUIRE(CHECK_COLUMN(result, 0, {1, 2, 3}));  
    REQUIRE(CHECK_COLUMN(result, 1, {2, 3, 4}));  
    REQUIRE(CHECK_COLUMN(result, 2, {1, 2, Value()}));  
    REQUIRE(CHECK_COLUMN(result, 3, {10, 20, Value()}));  
}
```

# Continuous Benchmarking

- Result verification and performance testing
- First correct, then fast
- Benchmarks
  - Microbenchmarks
  - TPC-H (complete)
  - TPC-DS (73/103 queries run)
  - TPC-E (only data generation)
- All use *in-process data generation*

```
DuckDB db(nullptr);  
DuckDBConnection con(db);  
tpch::dbgen(1, db);
```

```

DUCKDB_BENCHMARK(RangeJoin, "[micro]")
virtual void Load(DuckDBBenchmarkState *state) {
    // fixed seed random numbers
    std::uniform_int_distribution<> distribution(1, 10000);
    std::mt19937 gen;
    gen.seed(42);


    state->conn.Query("CREATE TABLE integers(i INTEGER, j INTEGER);");
    auto appender = state->conn.GetAppender("integers");
    // insert the elements into the database
    for (size_t i = 0; i < RANGEJOIN_COUNT; i++) {
        appender->begin_append_row();
        appender->append_int(distribution(gen));
        appender->append_int(distribution(gen));
        appender->end_append_row();
    }
    state->conn.DestroyAppender();
}

virtual std::string GetQuery() {
    return "SELECT * FROM integers a, integers b WHERE (a.i / 1000) > b.j;";
}

virtual std::string VerifyResult(DuckDBResult *result) {
    if (!result->GetSuccess()) {
        return result->GetErrorMessage();
    }
    return std::string();
}
FINISH_BENCHMARK(RangeJoin)

```

# Git revisions



| [micro]                | 0870             | 155a             | 9a39             | 39f3            | cb85            | be19            |
|------------------------|------------------|------------------|------------------|-----------------|-----------------|-----------------|
| Multiplication         | 0.11<br>[L/O/E]  | 0.11<br>[L/O/E]  | 0.11<br>[L/O/E]  | 0.11<br>[L/O/E] | 0.11<br>[L/O/E] | 0.11<br>[L/O/E] |
| OrderBySingleColumn    | 0.10<br>[L/O/E]  | 0.10<br>[L/O/E]  | 0.10<br>[L/O/E]  | 0.10<br>[L/O/E] | 0.10<br>[L/O/E] | 0.10<br>[L/O/E] |
| RangeJoin              | 21.66<br>[L/O/E] | 21.49<br>[L/O/E] | 21.50<br>[L/O/E] | 7.17<br>[L/O/E] | 7.29<br>[L/O/E] | 7.30<br>[L/O/E] |
| SimpleGroupByAggregate | 0.29<br>[L/O/E]  | 0.29<br>[L/O/E]  | 0.29<br>[L/O/E]  | 0.28<br>[L/O/E] | 0.29<br>[L/O/E] | 0.28<br>[L/O/E] |



Catches regressions

<http://www.duckdb.org/benchmarking/>

# Agenda

- Motivation
- Design
- Implementation
- Testing
- **Next Steps**



# Next Steps

- Help is appreciated, ask us for repo access and send a PR
  - Physical storage and buffer manager (Only WAL at the moment)
  - Foreign keys
  - Join ordering (Idea: use sampling) and more optimiser rules
  - Prepared statements & query cache
  - More types (real decimal, blob, timestamp)
  - More SQL features (`PARTITION` etc.)
  - Intra-query parallelism